

BSc (Hons) Artificial Intelligence and Data Science

Module: M2606 - Data Engineering

Individual Coursework Report

Module Leader: Mr. Mohamed Ayoob

RGU Student ID : 2410213

IIT Student ID : 20233168

Student Name :LOGANATHAN THUSHARKANTH

Contents

QUESTION 1 – OpenWeatherMap ETL TASK	5
1. Introduction	5
2. Setup & Configuration	5
2.1 Platform Selection: AWS EC2 (Ubuntu 22.04 LTS)	5
2.2 System Setup and Tool Installation	5
2.3 Creating an S3 Bucket and Configuring IAM Role	7
2.4 Setup for OpenWeatherMap API	7
3. Building the ETL Pipeline	7
3.1 API Availability Check – check_api	7
3.2 Extract – extract_task	8
3.3 Transform – transform_task	8
3.4 Load – load_task	9
4. Automation and Scheduling	9
5. Evaluation and Testing	10
5.1 Testing Summary	10
5.2 Challenges Faced & Solutions Applied	11
5.3 future improvements	12
6. Conclusion	13
7. Video Link	13
QUESTION 2 – SUPERMART	13
1. Identify Dimension Tables:	13
1.1 Store_Table	13
1.2 Product_Table	14
1.3 Customer_Table	14
1.4 Time_Table	15
1.5 Supplier_Table	15
2. Design of the Fact Table	16
3. Aggregate Tables	17
4. Business Questions	18
5. Diagram:	19

Question 3 - Security and Access Framework for	21
HealthAnalytics Inc.	21
1. Introduction	21
2. Authentication vs. Authorization	21
2.1 Authentication: Multi-Factor Authentication (MFA)	21
2.2 Authorization: Attribute-Based Access Control (ABAC) and Role-Based Access Control (RBAC)	21
2.3 RBAC Matrix	22
3. Data Encryption	22
3.1 Data Encryption at Rest: AES-256.....	22
3.2 Data Encryption in Transit: TLS 1.3.....	22
4. Access Control.....	23
4.1 Column-Level Security (CLS):	23
4.2 Dynamic Data Masking (DDM):	23
5. Compliance & Auditing	24
5.1 Audit Trails	24
5.2 Data Retention Policies.....	24
5.3 Real-Time Monitoring	25
6. Third-Party Integration	25
6.1 API Gateway	25
6.2 Token-Based Access	25
6.3 Rate Limiting	25
6.4 Encryption & Secure Communication	26
7. Architecture flowchart.....	26
8. Conclusion	28
References	29

Table of figures

Figure 1	5
Figure 2	5
Figure 3	5
Figure 4	6
Figure 5	6
Figure 6	6
Figure 7	7
Figure 8	7
Figure 9	8
Figure 10	8
Figure 11	9

<i>Figure 12</i>	9
<i>Figure 13</i>	10
<i>Figure 14</i>	10
<i>Figure 15</i>	10
<i>Figure 16</i>	11
<i>Figure 17</i>	12
<i>Figure 18</i>	12
<i>Figure 19: Star scheme</i>	19
<i>Figure 20: Snowflake Schema</i>	20
<i>Figure 21: Architecture flowchart</i>	27

QUESTION 1 – OpenWeatherMap ETL TASK

1. Introduction

This project builds an automated **ETL** pipeline using **Python, Apache Airflow, and AWS S3** to fetch real-time weather data from the **OpenWeatherMap API**, convert it into a structured JSON object, and store it in **AWS S3** as part of a data lake. Apache Airflow orchestrates the pipeline to run every hour, offering smooth data flow and guaranteed scheduling. It also provides hands-on training in basic data engineering concepts such as ETL development, cloud storage, workflow automation, and API integration, providing practical exposure to cloud platforms and tools.

2. Setup & Configuration

2.1 Platform Selection: AWS EC2 (Ubuntu 22.04 LTS)

I decided to use an **AWS EC2** instance with **Ubuntu 22.04 LTS** (HVM, SSD Volume Type) for this project due to compatibility issues with Ubuntu 24.04 and Apache Airflow. Ubuntu 22.04 LTS provides a stable and highly supported environment to run Airflow, making it more suitable for production-level implementation. On booting the instance, necessary ports were configured — **SSH, HTTP**, and a dedicated **TCP port** for the Airflow web UI — to enable proper access and communication. (AWS, n.d.)

2.2 System Setup and Tool Installation

After launching the EC2 instance, the system was updated to include all the dependencies and security patches updated in order to have a good foundation to install and run Apache Airflow.

```
# Update system
sudo apt update && sudo apt upgrade -y
```

Figure 1

Next, Python 3, pip, and virtualenv were installed to have a good Python environment in charge. Python 3 is the new standard of modern data engineering pipelines and is required for a good Airflow setup.

```
# Install Python and virtual env
sudo apt install python3-pip python3-venv -y
```

Figure 2

created a **Python virtual environment** (airflow_env) and activated it

```
# Create virtual environment
python3 -m venv airflow_env
source airflow_env/bin/activate
```

Figure 3

set specific versions for Apache Airflow and install Apache Airflow (Documentation, n.d.)

```
# Set versions
AIRFLOW_VERSION=2.8.1
PYTHON_VERSION=$(python --version | cut -d " " -f 2 | cut -d "." -f 1,2)
CONSTRAINT_URL="https://raw.githubusercontent.com/apache/airflow/constraints-${AIRFLOW_VERSION}/c

# Install Airflow
pip install "apache-airflow==${AIRFLOW_VERSION}" --constraint "${CONSTRAINT_URL}"
```

Figure 4

Airflow Configuration and Database Initialization

```
export AIRFLOW_HOME=~/.airflow
airflow db init
```

Figure 5

Creating an Admin User for Airflow

```
airflow users create \
  --username admin \
  --firstname yourname \
  --lastname yourname \
  --role Admin \
  --email your@email.com \
  --password admin
```

Figure 6

Starting Airflow Services (Webserver and Scheduler)

```
# Start Airflow services
source airflow_env/bin/activate
airflow webserver --port 8080
source airflow_env/bin/activate
airflow scheduler
```

Figure 7

tested the Airflow setup by accessing the **Airflow UI** at the following URL

```
http://<your-ec2-public-ip>:8080
```

Figure 8

2.3 Creating an S3 Bucket and Configuring IAM Role

I created an **S3 bucket** (my-data-engineering-coursework) to store the weather data. **AWS S3** (Services, n.d.) was utilized for its scalability and reliability as a data lake storage. To access securely from the EC2 instance that is running Airflow, an IAM role was created and attached and assigned EC2Full Access and S3Full Access permissions. (Services., n.d.)

2.4 Setup for OpenWeatherMap API

I registered on the OpenWeatherMap platform to obtain the API key, which is required to make authenticated requests to fetch real-time weather data (OpenWeather, n.d.)

3. Building the ETL Pipeline

This Airflow DAG runs **hourly** and performs 4 steps: **Check** → **Extract** → **Transform** → **Load**. Here's what each task does:

3.1 API Availability Check – check_api

The HttpSensor is used to ping the OpenWeatherMap API before the ETL process can be started. The **DAG executes** (dags, n.d.) only when the API is accessible and the response contains the keyword **"main."** This approach prevents failures of pipelines by making the API operational before it goes ahead to extract anything.

```

check_api = HttpSensor(
    task_id="check_api_availability",
    http_conn_id="openweather_api", # Defined in Airflow UI
    endpoint=f"data/2.5/weather?q={CITY}&appid={API_KEY}",
    method="GET",
    response_check=lambda response: "main" in response.text,
    poke_interval=10,
    timeout=60,
    dag=dag
)

```

Figure 9

3.2 Extract – extract_task

This task **fetches the current weather data for Portland** using the **OpenWeatherMap API** and saves the raw response as a **timestamped JSON file** in the /tmp/ directory. It uses a **GET request** to retrieve the data and stores it in a file named like `portland_weather_raw_<timestamp>.json`. This step is crucial because it **captures real-time weather data** and maintains a **versioned record** that can be used for tracking or auditing purposes.

```

def extract(**kwargs):
    execution_date = kwargs['ts_nodash']
    url = f"https://api.openweathermap.org/data/2.5/weather?q={CITY}&appid={API_KEY}"
    response = requests.get(url)
    data = response.json()
    raw_filename = f"/tmp/{CITY.lower()}_weather_raw_{execution_date}.json"
    with open(raw_filename, "w") as f:
        json.dump(data, f)

```

Figure 10

3.3 Transform – transform_task

This task reads the raw **JSON** file and extracts only the key fields such as city, timestamp, temperature, humidity, weather description, and wind speed. The cleaned and formatted data is then output to a new JSON file in the /tmp/ directory. This makes the data well formatted and ready for analysis or storage, without redundant or noisy data.


```
def transform(**kwargs):
    execution_date = kwargs['ts_nodash']
    raw_filepath = f"/tmp/{CITY.lower()}_weather_raw_{execution_date}.json"
    with open(raw_filepath) as f:
        data = json.load(f)

    transformed = {
        "city": data.get("name"),
        "timestamp": data.get("dt"),
        "temperature": data.get("main", {}).get("temp"),
        "humidity": data.get("main", {}).get("humidity"),
        "weather": data.get("weather", [{}])[0].get("description"),
        "wind_speed": data.get("wind", {}).get("speed")
    }

    transformed_filepath = f"/tmp/{CITY.lower()}_weather_transformed_{execution_date}.json"
    with open(transformed_filepath, "w") as f:
        json.dump(transformed, f)
```

Figure 11

3.4 Load – load_task

This task **uploads the transformed data (in JSON format) to an S3 bucket** using the **boto3** library. The data is stored under the path `weather-data/portland_weather_<timestamp>.json`. This step ensures that the **processed weather data is securely stored in the cloud**, making it easily accessible for **future analysis, reporting, or retrieval**.

```
def load(**kwargs):
    execution_date = kwargs['ts_nodash']
    transformed_filepath = f"/tmp/{CITY.lower()}_weather_transformed_{execution_date}.json"
    s3_key = f"weather-data/{CITY.lower()}_weather_{execution_date}.json"

    s3 = boto3.client("s3")
    s3.upload_file(transformed_filepath, S3_BUCKET, s3_key)
```

Figure 12

4. Automation and Scheduling

4.1 Schedule Interval:

The DAG is set to run automatically every hour with **@hourly**, which means the tasks will execute once every hour, ensuring that the data is always up-to-date.

```
dag = DAG(
    dag_id="weather_etl_dag",
    default_args=default_args,
    start_date=datetime(2024, 1, 1),
    schedule_interval="@hourly",
    catchup=False
)
```

Figure 13

4.2 Airflow Retry Logic:

Airflow takes care of retrying failed tasks by setting the `retries` parameter in the **default_args**. If a task fails, it will automatically **retry once after a 2-minute delay**, preventing manual intervention.

```
default_args = {
    "owner": "airflow",
    "retries": 1,
    "retry_delay": timedelta(minutes=2)
}
```

Figure 14

4.3 Logging Task Status:

Airflow automatically logs the status of each task, so you can monitor and track the progress of the ETL pipeline, ensuring that you are aware of any issues during execution.

4.4 Execution Context:

The **`provide_context=True`** argument in the **PythonOperator** allows access to execution metadata like **timestamps (`ts_nodash`)**. This is used to create timestamped filenames for raw and transformed data files, ensuring each file is uniquely identifiable.

```
extract_task = PythonOperator(
    task_id="extract_weather_data",
    python_callable=extract,
    dag=dag,
    provide_context=True # Access to ts_nodash, etc.
)
```

Figure 15

5. Evaluation and Testing

5.1 Testing Summary

1. Manual Triggering and Scheduled Execution:

The DAG was manually triggered from the Airflow UI, which allowed immediate testing of task execution. This ensured that the DAG could execute as expected when triggered manually.

The DAG was also set to execute on schedule by utilizing **@hourly**, meaning that it automatically triggers every hour. This automated triggering ensured that the scheduling process was

working well, executing tasks on expected intervals.

2. File Saving in /tmp/ Directory:

After each DAG run, raw and transformed JSON files were being saved in the /tmp/ directory. This indicated that the extract and transform tasks were being executed successfully, and the system was correctly handling the file creation after each task. The raw data was being extracted, and the transformation logic was being applied correctly, resulting in valid transformed files being saved locally for further processing.

3. S3 Upload Failure and Solution:

The upload to **AWS S3** failed in the load task initially. This was traced by viewing the Airflow logs and console output, which provided valuable details on what had failed. The error occurred because there was a Boto3 initialization error.

After some round of troubleshooting, including reinstalling Boto3 in the Airflow virtual environment and reviewing the logs, the issue was traced and fixed. Once the issue was fixed, the upload task could upload the transformed files to the target S3 bucket.

Manual verification was carried out by checking the S3 bucket through the AWS S3 Console, where the uploaded files were confirmed to be stored correctly. This final verification step ensured that the data had been migrated and made accessible in the cloud.




Figure 16

5.2 Challenges Faced & Solutions Applied

1. Boto3 Module Error

Airflow failed to load the task due to an import error with boto3. This was due to the fact that the boto3 module wasn't properly installed within the Airflow environment (venv). This caused the task to fail when attempting to push the converted file onto S3.

To fix this, I re-installed boto3 in the Airflow virtual environment (venv) and ensured the installation. Once that was done, the upload task too worked as expected, and the file uploaded successfully to

S3. The issue was fixed by ensuring the boto3 module had been installed correctly, and the upload process ran without issues.

```
pip install boto3
```

Figure 17

2. File Overwrite Issue Due to Same Filenames

To start with, I used `kwargs['ds']`, which retrieved only the date (e.g., 20250420) and hence provided the same filename when the DAG was run every time. This caused S3 to overwrite the previously uploaded file because the names were the same.

To fix this, I switched to utilizing `kwargs['ts_nodash']`, which includes the timestamp (like 202504201200). This adjustment ensured that each DAG run produced a unique file from a timestamped filename so that no S3 files were clobbered. As a result, all files had a distinct name and uploaded successfully without overwriting previous ones.

```
execution_date = kwargs['ts_nodash']
transformed_filepath = f"/tmp/{CITY.lower()}_weather_transformed_{execution_date}.json"
```

Figure 18

5.3 future improvements

1. Use S3Hook Instead of Raw Boto3

S3Hook is the recommended way of working with S3 in Airflow. It simplifies hooking into AWS, adds additional security through the automated management of credentials, and maintains the code simpler to write by bringing Airflow's connection management into it. Switching to S3Hook will secure and streamline your file uploads.

2. Use Logging Module Instead of print()

Airflow works best when logs are processed through the logging module. Using logging instead of print statements means that logs can be captured in the Airflow UI, so you can simply view task status and debug problems. This change will improve the robustness of your error handling and make the workflow more professional.

3. Store API Key in Airflow Variables

Hardcoding API keys is insecure. Storing API keys as Airflow Variables allows you to centralize and safely store sensitive information. This reduces the likelihood of exposing your credentials and adds a layer of security to your project as a whole.

4. Utilize XCom to Share File Paths Between Tasks

Sharing data, like file paths, between tasks with the help of XCom makes your workflow more modular and easier to manage. It ensures your tasks are decoupled, reducing dependencies and maximizing your DAG's flexibility. It also helps with task-specific data management and enhances maintainability for your code.

5. Add a FileSensor Before Transform

To prevent errors during data conversion, the presence of a File Sensor will ensure that the raw file exists before any conversions are attempted. This operation will prevent tasks from running prematurely, ensuring that data is available for processing when needed. It will also reduce the likelihood of errors from missing or incomplete files.

6. Conclusion

This project showcases successfully constructing and automating an end-to-end real-time data pipeline using **Airflow, Python, and AWS**, showcasing competence in API integration, data processing, and cloud storage using AWS S3. The pipeline performs data extraction, transformation, and loading (ETL) very efficiently with negligible human intervention. Key takeaways are seamless integration with AWS services for safe cloud storage, automatic handling of weather data, and Airflow for processing intricate workflows. The project can be extended further by incorporating additional cities or weather conditions, storing data in databases to query, and integrating with dashboarding platforms like **Amazon QuickSight** or **Grafana** for real-time analytics. Overall, this project gives a strong foundation for scalable, self-automating data pipelines that possess the future ability to grow and adapt.

7. Video Link

<https://drive.google.com/file/d/1cNqBaZgcviXH04R707QccRQ1BtKFB6UM/view?usp=sharing>

QUESTION 2 – SUPERMART

1. Identify Dimension Tables:

1. Store_Table

Attribute	Description
store_id (PK)	Unique identifier for each store.
store_name	Name of the store.
city	City where the store is located.
state	State where the store is located.

country	Country where the store is located.
region	Business/administrative region.

Role:

This dimension table contains the geographical information of each store. It supports analysis by store location (city, state, country) and can be utilized for the comparison of the sales performance of different regions.

2. Product_Table

Attribute	Description
product_id (PK)	Unique identifier for each product.
product_name	Name of the product.
category	Category of the product (e.g., electronics, clothing).
brand	Brand of the product.
unit_price	Price of the product per unit.
supplier_id / supplier_Name	ID of the supplier for the product.

Role:

This table contains product-specific information, which helps in product-wise sales performance analysis, price, inventory management, and supplier details. It enables detailed analysis of sales by supplier, brand, and category.

3. Customer_Table

Attribute	Description
customer_id (PK)	Unique identifier for each customer.
customer_name	Full name of the customer.
address	address of the customer.

membership_level	Loyalty program level (bronze, silver, gold).
------------------	---

Role:

This table contains customer profile data, which is important for loyalty analysis, target marketing, and customer segmentation. SuperMart can analyze customer information by monitoring loyalty schemes by membership level, or location to send personalized campaigns.

4. Time_Table

Attribute	Description
time_id (PK)	Unique identifier for each time period
date	Actual date of the transaction.
day	Day of the month (1–31).
month	Month (e.g., January, February, etc.).
quarter	Quarter of the year (Q1–Q4).
year	Year (e.g., 2025).
weekday	Name of the day (e.g., Monday, Tuesday, etc.).
is_weekend	Indicates whether the day is a weekend (Yes/No).

Role:

The Time dimension table supports the examination of sales information based on different time levels such as daily, monthly, quarterly, or annually. It allows for the time-related study of sales, the seasonal trend, and the sales pattern analysis based on different days of the week or by quarter.

5. Supplier_Table

Attribute	Description
supplier_id (PK)	Unique identifier for each supplier.

supplier_name	Name of the supplier.
supplier_address	Address of the supplier.
contact_number	Contact information for the supplier.
email_address	Supplier's email address.

2. Design of the Fact Table

The fact table contains transactional-level data that has measurable facts about sales. It is the fact table in the star schema and joins with dimension tables through foreign keys. It supports analysis of sales performance by different time period, stores, products, and customers.

Fact Table Name: Sales_Transactions_Table

Attributes of the Fact Table:

Attribute	Description
transaction_id	Primary key. Unique identifier for each sales transaction.
time_id	Foreign key linked to the Time dimension. Represents the transaction date.
store_id	Foreign key linked to the Store dimension. Indicates where the sale happened.
product_id	Foreign key linked to the Product dimension. Identifies the sold product.
customer_id	Foreign key linked to the Customer dimension. Identifies who made the purchase.
quantity_sold	Numeric fact. Total units sold in the transaction.
total_sales_amount	Numeric fact. Total amount generated from the sale (quantity × unit price).

Foreign Key Relationships:

- store_id → Store dimension
- product_id → Product dimension
- customer_id → Customer dimension

- time_id → Time dimension
-

Note: The Supplier dimension is not directly linked to the fact table but is accessible through the Product dimension, since each product has a supplier_id

Type of Facts and Their Significance:

Fact Column	Type	Explanation
quantity_sold	Additive	Can be summed across any dimension (e.g., total quantity sold per month).
total_sales_amount	Additive	Can be summed across all dimensions to analyze revenue (e.g., by store or region).

Additive facts are used because they support **aggregation** over any combination of dimensions, making them essential for business reporting and analysis.

3. Aggregate Tables

To enhance query performance, two aggregate tables may be established to precompute frequent business measures

1. Monthly_Sales_Aggregate_Table

Purpose:

Rolls up total sales and quantity sold for each store and product by month. This allows for prompt response to such questions as:

"What are the total sales for each store this month?"

Attributes:

- store_id (FK)
- product_id (FK)
- month
- year
- total_sales_amount (sum of sales)
- total_quantity_sold (sum of quantity)

Relation:

- store_id and product_id link to the Store_Table and Product_Table.
- month and year link to the Time_Table.
- Speeds up monthly sales reporting by precomputing totals.

2. Yearly_Sales_Aggregate_Table

Purpose:

Precomputes total sales and quantity sold for each store and product per year. This helps answer questions like:

- “Which store had the highest sales in 2025?”

Attributes:

- store_id (FK)
- product_id (FK)
- year
- total_sales_amount (sum of sales)
- total_quantity_sold (sum of quantity)

Relation:

- store_id and product_id link to Store_Table and Product_Table.
- year links to the Time_Table.
- Speeds up yearly sales reporting by storing precomputed data.

Benefits:

- Improved Performance: Reduces the need to re-calculate totals from the large fact table.
- Faster Reporting: Allows quick access to summarized sales data by month and year.
- Reduced Load: Minimizes database strain by storing aggregate data in precomputed tables.

4. Business Questions

1. Which store had the highest sales last month?

Uses **Monthly_Sales_Aggregate_Table** to compare **total_sales_amount** by **store_id**.

This question identifies the top-performing store by comparing total monthly sales. It uses pre-aggregated data to give instant insight into store performance.

2. What are the top-selling products by quantity in each region this year?

Uses **Yearly_Sales_Aggregate_Table** + joins with **Store_Table** to group by region and product.

This question identifies the top-selling products in each region by looking at yearly sales quantity. It combines yearly data with store information for regional grouping.

3. Which product categories are generating the most revenue over the year?

Uses **Yearly_Sales_Aggregate_Table** + join with **Product_Table** to group by category and sum revenue.

This question identifies which product categories bring in the most revenue by grouping by category and aggregating total sales over the year.

5. Diagram:

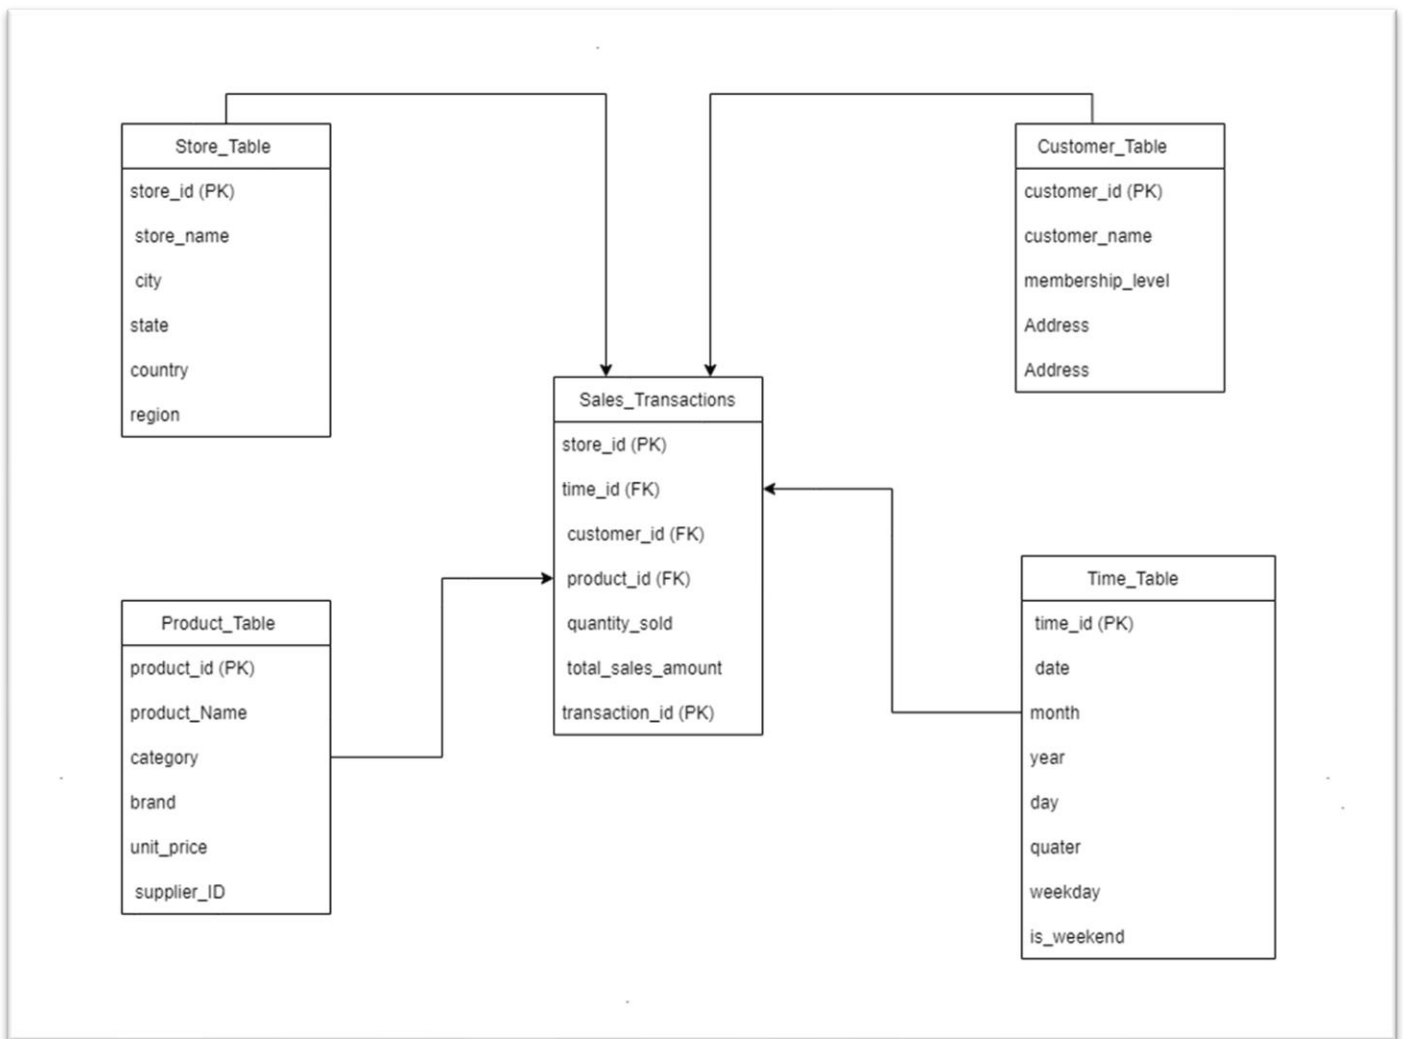


Figure 19: Star scheme

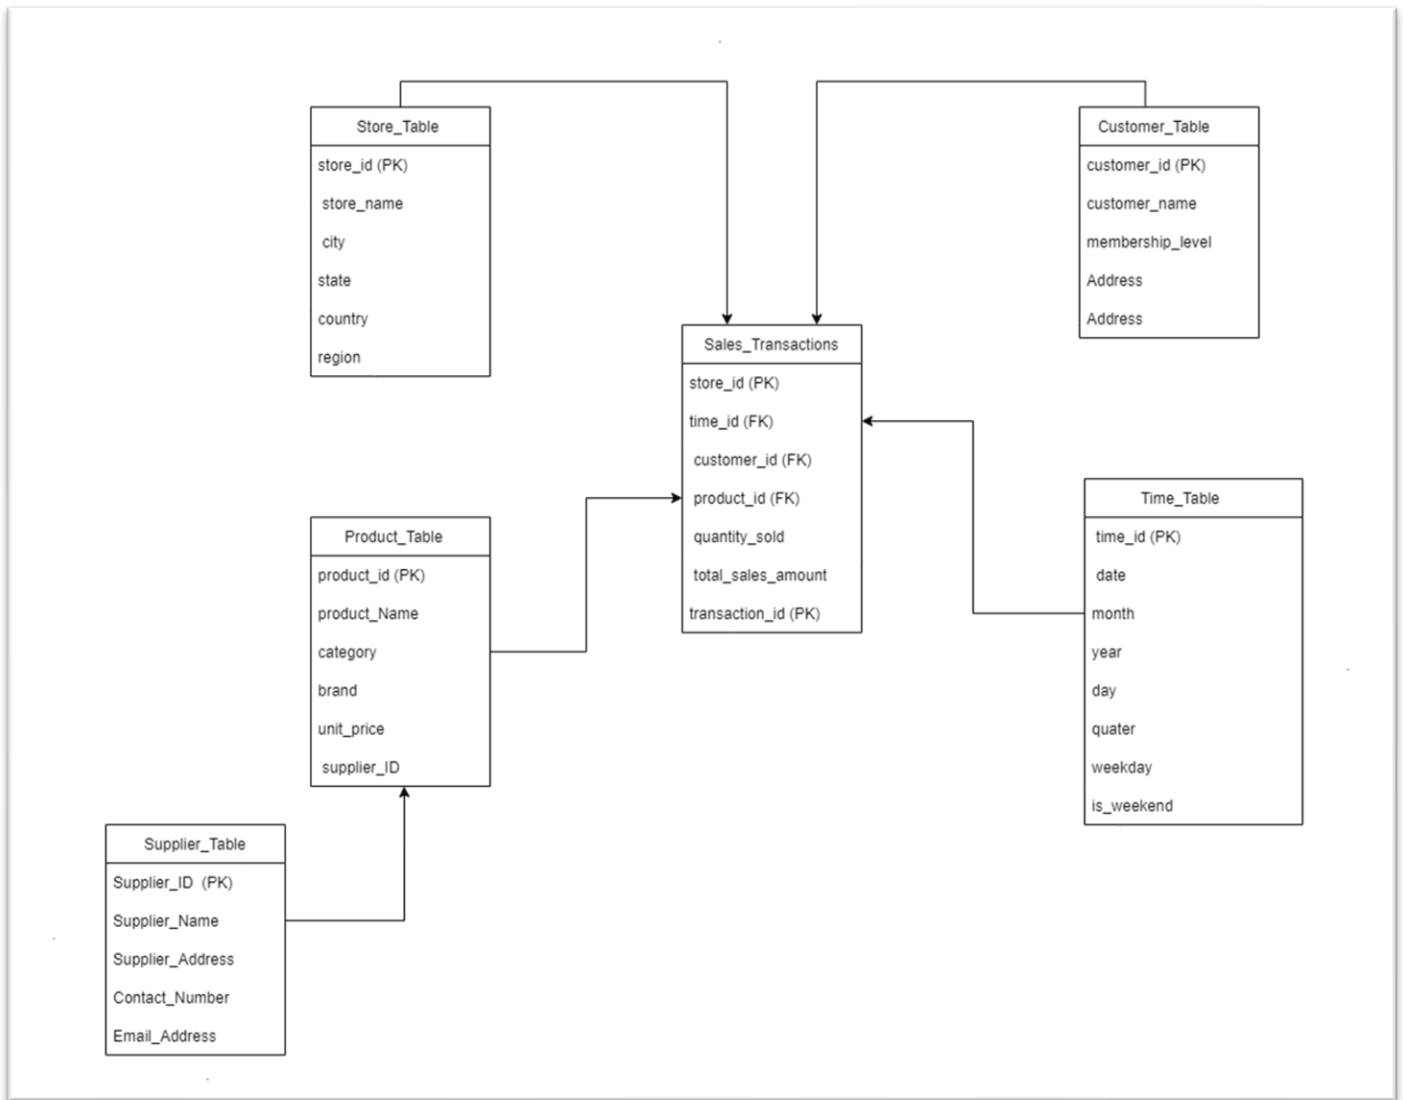


Figure 20: Snowflake Schema

Note on Schema Diagrams:

We have included **two different schema diagrams** for clarity and completeness:

1. **Star Schema** – This is the simplified version where the **Supplier Table is excluded**, as it is not directly connected to the Fact Table.
2. **Snowflake Schema** – This version includes the **Supplier Table**, connected through the Product Table. This normalization reflects real-world relationships more accurately.

While adding the Supplier as a separate dimension is **optional** in this case, it's a **good design practice** as it improves data organization and allows supplier-level analysis when needed.

Question 3 - Security and Access Framework for HealthAnalytics Inc.

1. Introduction

Protected Health Information (PHI) and **Personally Identifiable Information (PII)** are among the sensitive health care data that Health Analytics Inc. is currently moving to a cloud-based data lakehouse. In order to support medical research and large-scale analytics, which will spur innovation and enhance patient outcomes, this migration is essential. However, because health care data is sensitive, it is essential to make sure that all data is protected in accordance with strict laws like the General Data Protection Regulation (GDPR) and the Health Insurance Portability and Accountability Act (HIPAA).

The scalable, secure architecture described in this report uses Amazon Web Services (AWS) to protect sensitive data during the migration process. In addition to improving data security, the architecture will guarantee that authorized users can effectively access and analyze data for research and clinical decision-making, while minimizing the risk of unauthorized access or data breaches.

2. Authentication vs. Authorization

2.1 Authentication: Multi-Factor Authentication (MFA)

Multi-Factor Authentication (MFA) will be used to guarantee safe access to the cloud-based data lakehouse. Before being allowed access, MFA requires users to present two or more forms of verification, including:

- Something you are familiar with (PIN or password).
- Something you own (OTP sent by email or mobile).

By lowering the possibility of unwanted access, this technique greatly improves the security posture, particularly for users handling private medical information. In accordance with security best practices and compliance requirements (**HIPAA and GDPR**), MFA will be enforced using both Azure Active Directory (AD) and AWS Identity and Access Management (IAM).

2.2 Authorization: Attribute-Based Access Control (ABAC) and Role-Based Access Control (RBAC)

Role-Based Access Control (RBAC) and **Attribute-Based Access Control (ABAC)** will be used in tandem for granular authorization.

Role-Based Access Control (RBAC): RBAC is in line with the HIPAA "**minimum necessary**" rule and the "**least privilege**" principle by limiting access to resources based on the user's assigned

rele.

Attribute-Based Access Control (ABAC): ABAC gives you more precise control over resource access by enforcing policies dynamically based on user attributes (like department and shift timing).

Only authorized personnel can access sensitive patient data thanks to these models, which grant specific access based on roles and contextual factors.

2.3 RBAC Matrix

The following table outlines the access permissions for different roles within the system:

Role	Permission
Doctors	Full access to assigned patient records, including editing medical history.
Data Scientists	Access to anonymized, aggregated data for analysis, no direct access to identifiable patient data.
Auditors	Read-only access to logs, non-PII fields, and compliance reports.

ABAC Example

- A **Data Scientist** can access data related only to their assigned research domain (e.g., oncology).
- A **Doctor** can view patient records only during their shift.

Tools like Apache Ranger and AWS Lake Formation, which interface with IAM to guarantee that the appropriate level of access is applied based on the user's role and attributes, will be used to enforce these policies. For added security, sensitive fields like **Social Security Numbers (SSN)** or Insurance IDs will be hidden from unauthorized users using Dynamic Data Masking.

3. Data Encryption

To ensure the security of sensitive healthcare data, both at rest and in transit, the following encryption methods are recommended:

3.1 Data Encryption at Rest: AES-256

The standard for encrypting sensitive data is AES-256 (Advanced Encryption Standard with a 256-bit key), which makes sure that the confidentiality of stored data is not jeopardized by unauthorized access.

Reasoning:

- AES-256 offers robust security and complies with GDPR and HIPAA regulations.
- Cloud service providers like AWS and Azure support it extensively.
- Encryption keys can be managed by cloud services like Azure Key Vault or AWS KMS.

3.2 Data Encryption in Transit: TLS 1.3

The most recent version, TLS 1.3 (Transport Layer Security), encrypts data while it is in transit to enable secure network communication.

Reasoning:

- Interception and man-in-the-middle attacks are prevented by TLS 1.3.
- Forward secrecy is supported, protecting previous communications even in the event that future keys are stolen.
- ensures that data transmission complies with GDPR and HIPAA.

Health Analytics Inc. will protect sensitive healthcare data while adhering to regulatory requirements by deploying TLS 1.3 for data in transit and AES-256 for data at rest.

4. Access Control

The following access control procedures will be used to limit access to private information, including Social Security Numbers (SSNs):

4.1 Column-Level Security (CLS):

Goal: Make sure that only authorized users can access sensitive fields, such as SSNs.

Execution:

- SSNs for the patients they are assigned will be fully accessible to doctors.
- Only anonymized data with masked SSNs will be available to data scientists.
- Only non-sensitive data with masked SSNs will be visible to auditors.

AWS Service: To efficiently manage column-level permissions, AWS Lake Formation will be used.

4.2 Dynamic Data Masking (DDM):

The goal is to conceal private information from users who lack the required authorization.

Execution: Unauthorized users will see SSNs in a masked format (such as "XXX-XX-XXXX").

AWS Service: Depending on user roles, dynamic data masking will be applied using AWS Redshift.

By limiting access to sensitive data to authorized personnel, this policy upholds HIPAA and GDPR compliance.

5. Compliance & Auditing

Protecting sensitive healthcare data requires adherence to legal frameworks like HIPAA and GDPR. HealthAnalytics Inc. must put in place reliable auditing and monitoring systems in addition to data retention guidelines in order to satisfy these requirements. To guarantee compliance, the following actions will be taken.

5.1 Audit Trails

Goal: To guarantee accountability and transparency, keep thorough records of all accesses to and changes made to sensitive healthcare data.

Execution:

- Turn on thorough logging for all queries, updates, and accesses involving private patient information.
- Logs will document user activities, including who accessed the data, what was accessed, and when.

AWS tool:

All API calls are logged using **AWS CloudTrail**, which creates an unchangeable audit trail of user interactions with the data lakehouse environment.

5.2 Data Retention Policies

Goal: In accordance with HIPAA and GDPR, securely delete sensitive data after it is no longer needed and only keep it for as long as is necessary.

Execution:

- Policies for data retention will be established to automatically archive or remove data in accordance with retention requirements.
- The "Right to be Forgotten" under the GDPR will be put into effect, enabling patients to ask for their data to be deleted whenever they so choose.

AWS Tools:

Using **AWS S3 Lifecycle Policies**, patient data can be automatically archived or deleted according to predetermined retention periods.

5.3 Real-Time Monitoring

Goal: Prevent possible security breaches by immediately identifying and reacting to suspicious activity.

Execution:

- Keep an eye out for any attempts at illegal access, odd user behavior (such as logging in from odd places), or access beyond roles that have been assigned.
- Set off alerts in responses to questionable activity, guaranteeing prompt action.

AWS Tools:

- **AWS GuardDuty:** For ongoing security threat monitoring.
- Using **AWS CloudWatch**, administrators can create custom alerts based on security thresholds to be informed of unusual activity.
- **SIEM tools** (such as **IBM QRadar** and **Splunk**): For combining and examining log data in order to find any irregularities or security incidents.

HealthAnalytics Inc. will guarantee complete compliance with HIPAA and GDPR regulations by putting these auditing and compliance procedures into place. This will shield private medical information from unwanted access and guarantee that it is only kept for as long as is required.

6. Third-Party Integration

The following actions will be taken in order to safely provide a research partner with aggregated non-PII data:

6.1 API Gateway

The AWS API Gateway will be used to manage access to non-PII data. This preserves sensitive data while guaranteeing that only aggregated, non-PII data is available. Strict security controls for data exposure will be enforced by the API Gateway, which will act as the interface between the research partner and the data.

6.2 Token-Based Access

Token-based authentication will be used to make sure that only people with permission can access the data. Unique API tokens created using AWS Cognito or OAuth 2.0 will be given to research partners, giving them temporary, secure access to the data. Only verified users are allowed access thanks to this token-based system.

6.3 Rate Limiting

Rate limiting will be used to guard against abuse of the API and guarantee equitable use. For instance, external users who access the data via the API Gateway will be subject to a limit of 100 requests per minute. By doing this, you can keep the system operating smoothly and avoid overload or denial of

service attacks.

6.4 Encryption & Secure Communication

TLS 1.3 will be used to secure the data while it is in transit. Every API interaction will take place over HTTPS, guaranteeing that all correspondence between the data lakehouse and the research partner is encrypted and shielded from possible interception.

By utilizing AWS API Gateway, token-based access, rate limiting, and encryption, this approach guarantees that the research partner can safely access the combined, non-PII data without jeopardizing sensitive data.

7. Architecture flowchart

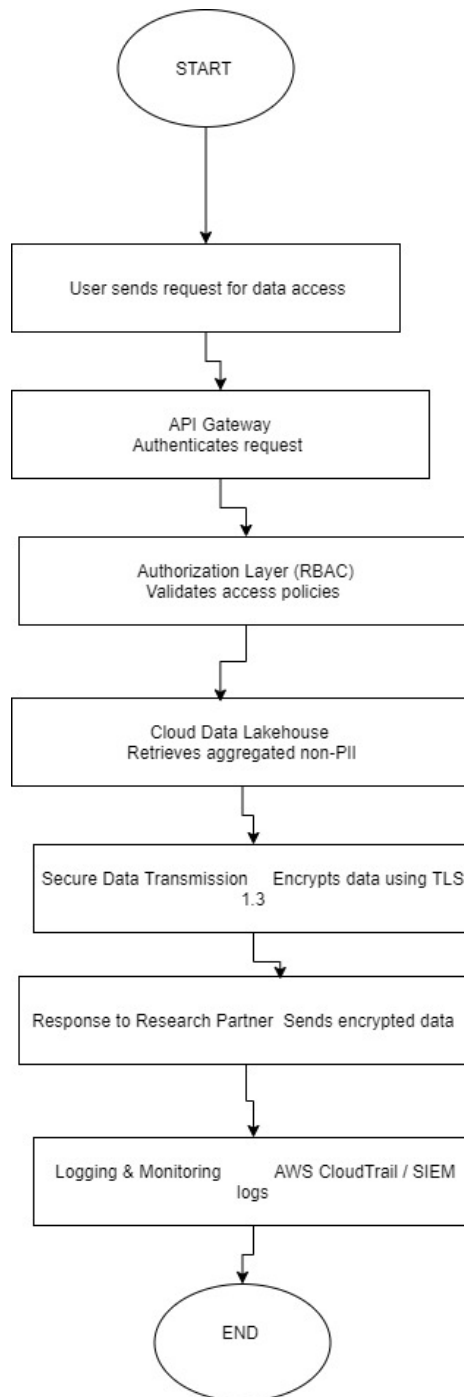


Figure 21: Architecture flowchart

8. Conclusion

In conclusion, HealthAnalytics Inc. must move sensitive medical data to a cloud-based data lakehouse in order to support research and enhance patient outcomes. To ensure that patient data is shielded from abuse, unauthorized access, and potential breaches, it is imperative that this migration be carried out in complete compliance with legal requirements such as **HIPAA and GDPR**.

HealthAnalytics Inc. can guarantee that only authorized personnel have access to sensitive data by putting in place a strong security and access framework that includes column-level security, dynamic data masking, encryption techniques like **AES-256** and **TLS 1.3**, Multi-Factor Authentication (**MFA**), Role-Based Access Control (**RBAC**), and Attribute-Based Access Control (**ABAC**). These actions enhance the organization's overall security posture in addition to meeting compliance requirements.

Audit trails, data retention guidelines, and real-time monitoring with **AWS CloudTrail**, **AWS GuardDuty**, and **SIEM** tools will all help to secure the data and guarantee that any questionable activity is promptly identified and addressed. Lastly, HealthAnalytics Inc. can securely share aggregated, non-PII data for research purposes without jeopardizing the integrity of sensitive information by utilizing a secure **API gateway**, token-based authentication, and rate limiting for third-party integrations.

In the end, this all-encompassing security and access framework will allow HealthAnalytics Inc. to function effectively in a safe, legal manner while cultivating an atmosphere that encourages creativity in medical research.

References

- AWS, n.d. **AWS-EC2 Instances**. [Online]
Available at: <https://eu-north-1.console.aws.amazon.com/ec2/home?region=eu-north-1#Instances:>
- dags, n.d. [Online]
Available at: <https://airflow.apache.org/docs/apache-airflow/stable/core-concepts/dags.html>
- Documentation, A. A., n.d. *Apache Airflow Documentation*. [Online]
Available at: <https://airflow.apache.org/docs/>
- OpenWeather, n.d. *OpenWeather API*. [Online]
Available at: <https://openweathermap.org/current>
- Services., A. W., n.d. *mazon Simple Storage Service Documentation*.. [Online]
Available at: <https://docs.aws.amazon.com/s3/>
- Services, A. W., n.d. *Getting started with Amazon S3*.. [Online]
Available at: <https://aws.amazon.com/getting-started/hands-on/>
- <https://www.youtube.com/watch?v=uhQ54Dgp6To>
- <https://youtu.be/5peQThvQmQk>
- Draw.io for the digrams -- <https://app.diagrams.net/>
- Current weather data -- <https://openweathermap.org/current>
- AWS-EC2 Instances. -- [Instances | EC2 | eu-north-1](#)
- [ETL Process in Data Warehouse | GeeksforGeeks](#)
- [Directed Acyclic Graph in Compiler Design \(with examples\) | GeeksforGeeks](#)
- [S3 buckets | S3 | eu-north-1](#)
- *Apache Airflow Documentation*..--<https://airflow.apache.org/docs/apache-airflow/stable/core-concepts/dags.html>
- [AWS Command Line Interface Documentation](#)
- [The Security Rule | HHS.gov](#)
- [General Data Protection Regulation \(GDPR\) – Legal Text](#)
- [Data protection in AWS Identity and Access Management - AWS Identity and Access Management](#)
- [Data Protection in Lake Formation - AWS Lake Formation](#)
- [Security design principles - Security Overview of Amazon API Gateway](#)
- [Data Protection and Privacy | AWS](#)

- [Guide to TLS Standards Compliance - SSL.com](#)
- [Compliance validation for AWS CloudTrail - AWS CloudTrail](#)
- [GDPR - Amazon Web Services \(AWS\)](#)
- [Controlling and Managing Access to a REST API in Amazon API Gateway](#)
- [How to Setup AWS Multi-Factor Authentication \(MFA\) and AWS Budget Alerts | Amazon Web Services](#)
- [Using Amazon GuardDuty to continuously monitor and secure your AWS AI workloads](#)
- Lecture notes and tutorials
- ChatGPT for Error Handling in Apache Airflow
- ChatGPT for Enhancing English Writing
- [airflow.hooks.S3_hook — Airflow Documentation](#)
- [python - How to use the s3 hook in airflow - Stack Overflow](#)
- [FileSensor — Airflow Documentation](#)
- [What are the design schemas of data modelling? | GeeksforGeeks](#)